

rxntoarb

v. 2.0

Christian Biscombe

7 September 2017

rxntoarb converts a human-readable system of chemical reactions into a format suitable for use with *arb* finite volume solver. Reading from one or more input files (with the suggested extension `.rxn`), **rxntoarb** automatically creates

- rate constant definitions (performing unit conversions where necessary),
- reaction rate expressions,
- source terms,
- transport equations for each chemical species (easily customisable by means of a template file),
- magnitude definitions.

Other features include easy switching between ODE and PDE output and the ability to include or exclude reactions based on regular expression matches.

rxntoarb offers several benefits:

- the reaction network is represented using intuitive syntax that is easier and faster to read and write,
- reactions may be easily added or removed without the user having to keep track of the effects of these changes on source terms,
- the potential for error is greatly reduced by eliminating the need for manual unit conversions and manual entry of reaction rates and source terms.

rxntoarb requires Ruby 1.9 or newer. Its output has been tested with *arb* versions 0.56 and 0.57.

Contents

1	Invocation and command-line options	2
2	File structure	3
2.1	.rxn file	3
2.2	Template file (rxntoarbc)	3
2.3	.arb output file	4
2.4	Main .arb file	5
3	.rxn file syntax	6
3.1	Comments	6
3.2	Headers	6
3.3	Species and region names	6
3.4	surface_region and volume_region statements	7
3.4.1	Lazy region identifiers @s and @v	7
3.5	Reactions	7
3.5.1	Irreversible reactions	8
3.5.2	Reversible reactions	8
3.5.3	Two-step reactions	8
3.5.4	Michaelis–Menten reactions	8
3.6	Kinetic parameters	9
3.6.1	Inheritance	9
3.7	Aliases	9
3.8	initial_species statement	10
3.9	include_only and exclude statements	10
4	Template file syntax	11
4.1	Substitution sequences	11
4.2	if_rxn statement	11
4.3	each blocks	12
5	Future development ideas	12
6	Copyright and licence	13

1 Invocation and command-line options

rxntoarbc is invoked as follows:

```
rxntoarbc [options_list] <rxn_file_list>,
```

where rxn_file_list specifies the path to at least one .rxn file and the optional options_list may include any of the following:

`-d|--debug` Write copious debugging output to STDERR. (Intended mainly as a developer option.)

`-i|--interactive` Prompt for confirmation before overwriting an existing `.arb` output file.

`-n|--none-centred` Activate ODE mode: generate ODEs (no spatial dependence) rather than PDEs. `rxntoarb` operates in PDE mode by default.

`-o|--outfile <output_file>` Use `output_file` as the `.arb` output file. This option is ignored if `rxn_file_list` contains more than one `.rxn` file.

`-t|--template <template_file>` Use `template_file` as the template file.

`-v|--version` Print version information.

2 File structure

Using `rxntoarb` with *arb* involves the use of four files.

2.1 .rxn file

`.rxn` files are input files read by `rxntoarb`. They contain a list of chemical reactions together with kinetic data. The syntax used in `.rxn` files is defined in § 3. An example `.rxn` file, reproduced in Listing 1, may be found under `examples/reactions.rxn` in the `rxntoarb` root directory. Vim syntax highlighting for `.rxn` files is provided by the file `rxn.vim` included in the `rxntoarb` root directory (type `:help usr_06` in Vim for instructions on how to set this up).

2.2 Template file (rxntoarbrc)

Template files store *arb* code that needs to be included for each chemical species defined in a `.rxn` file. The extended (non-*arb*) syntax used in template files is defined in § 4. `rxntoarb` uses the following hierarchy to determine the path of the template file:

- If `rxntoarb` is invoked with the `-t` flag, the specified template file is used.
- If a file with the name `rxntoarbrc` exists within the working directory from which `rxntoarb` is called, it is used as the template file.
- The default `rxntoarbrc` located in the `rxntoarb` root directory is used otherwise.

The default `rxntoarbrc` should provide an adequate starting point in most cases. However, the user may wish to alter the (simulation-specific) boundary conditions listed beneath the comment ‘`# set boundary conditions here`’ or to define new per-species `OUTPUT` variables. In this case, the default template should be copied to the working directory and modified there.

Listing 1 An example .rxn file (examples/reactions.rxn)

```
1  ! Blood plasma coagulation biochemistry
2  ! Ref: Jordan and Chaikof (2011), Biophys. J. 101: 276-286
3  ! File created by Christian Biscombe, 2017-08-19

5  # Fluid-phase reactions
6  volume_region domain
7  VIII@v {IIa@v}-> VIIIA@v; kcat=0.9 s-1, KM=0.20 uM
8  V@v {IIa@v}-> Va@v; kcat=0.23 s-1, KM=0.072 uM
9  VII@v {Xa@v}-> VIIa@v; kcat=15.2 s-1, KM=1.2 uM # CB corrected erroneous kcat in ref.
10 VII@v {IIa@v}-> VIIa@v; kcat=0.061 s-1, KM=2.7 uM
11 IIa@v + ATIII@v -> ; k=6.8d-3 uM-1 s-1
12 IXa@v + ATIII@v -> ; k=9.9d-6 uM-1 s-1 # CB corrected erroneous k in ref.
13 Xa@v + ATIII@v -> ; k=2.6d-3 uM-1 s-1 # reaction mistakenly omitted in ref.
14 Xa@v + TFPI@v <=> TFPI:Xa@v; ka=16 uM-1 s-1, kd=3.3d-4 s-1
15 Va@v {APC@v}-> ; kcat=0.58 s-1, KM=0.025 uM
16   VIIIA@v {APC@v}->
17 V@v {Xa@v}-> Va@v; kcat=0.046 s-1, KM=0.01 uM

19 # Reactions occurring on both TF patch and TM patch
20 surface_regions TF_patch, TM_patch
21 II@v {Va:Xa@s}-> IIa@v; kcat=33 s-1, KM=0.21 uM
22 X@v {VIIIA:IXa@s}-> Xa@v; kcat=20 s-1, KM=0.16 uM
23 Va@v + Xa@v <=> Va:Xa@s; ka=100 uM-1 s-1, kd=0.01 s-1 # units inconsistent!
24 VIIIA@v + IXa@v <=> VIIIA:IXa@s; ka=100 uM-1 s-1, kd=0.01 s-1 # units inconsistent!

26 # Reactions occurring on TF patch only
27 surface_region TF_patch
28 X@v {TF:VIIa@s}-> Xa@v; kcat=1.2 s-1, KM=0.45 uM
29 IX@v {TF:VIIa@s}-> IXa@v; kcat=0.34 s-1, KM=0.17 uM
30 VIIa@v + TF@s <=> TF:VIIa@s; ka=100 uM-1 s-1, kd=0.06 s-1
31 TF:VIIa@s + TFPI:Xa@v <=> TF:VIIa:TFPI:Xa@s; ka=10 uM-1 s-1, kd=1.1d-3 s-1

33 # Reactions occurring on TM patch only
34 surface_region TM_patch
35 PC@v {IIa:TM@s}-> APC@v; kcat=5.58 s-1, KM=0.7 uM
36 IIa@v + TM@s <=> IIa:TM@s; ka=100 uM-1 s-1, kd=0.01 s-1

38 initial_species [II, V, VII, VIIa, VIII, IX, X, PC, ATIII, TFPI]@domain
39 initial_species TF@TF_patch, TM@TM_patch
```

2.3 .arb output file

The output of `rxntoarb` is an .arb output file, which contains valid *arb* code metaprogrammed from the contents of the .rxn file. The .arb output file will have the same basename as the .rxn file unless overridden by the `-o` flag.

2.4 Main .arb file

The .arb output file generated by `rxntoarb` contains *arb* code describing chemical species and the reactions occurring between them, but it is not sufficient to run a complete *arb* simulation. Additional code related to mesh setup, solver options, equations governing fluid flow, time-stepping, etc. should be included in a separate .arb file, referred to as the main .arb file. It is this main .arb file (and not the .arb output file) that should be run through *arb*.

The main .arb file should have the following characteristics to interface cleanly with the .arb output file:

- An `INCLUDE_LOCAL` statement to include the .arb output file must be present.
- Any variables used in the template file but not created by `rxntoarb` must be defined.
- Replacements for any flags (strings delimited by double angle brackets `<<>>`) used in the template file must be defined. The default `rxntoarbrc` contains the following flags:
 - `<<speciesloworderadvection>>` (whether to use low-order advection for chemical species; should be either `"` or `#`)
 - `<<specieshighorderadvection>>` (whether to use high-order advection for chemical species; should be the opposite of `<<speciesloworderadvection>>`)
 - `<<zeroinitialvolumeconcentrationsflag>>` (whether volume species are initially present throughout the volume; should be either 0 or 1)
 - `<<calculatediffusivities>>` (whether species diffusivities should be calculated using a correlation based on molecular weight; should be either `"` or `#`)
- Species-specific physical parameters, currently only molecular weights or diffusivities, should be defined as `CONSTANTS`. Molecular weight variable names must follow the convention `<MW_speciesname>`. Diffusivity variable names must follow the convention `<D_speciesname>`.
- Initial concentrations of all species present initially should be defined as `CONSTANTS` *after* the .arb output file is included (because this file sets the initial concentrations of all species to zero by default). The initial concentration variable names must follow the convention `<c_speciesname@region_0>` for species located in volume regions and `<s_speciesname@region_0>` for species located on surface regions.

An example (skeleton) main .arb file may be found under `examples/main.arb` in the `rxntoarb` root directory.

3 .rxn file syntax

3.1 Comments

Comments begin with the `#` character and continue until the end of the line. Comments may appear anywhere in the .rxn file. Full-line comments (e.g. line 5 in Listing 1) are ignored by `rxntoarb`. Comments appearing after the parameter list of a reaction (e.g. line 9) are carried over into the .arb output file, where they appear after the `CONSTANT` expressions that define the relevant kinetic parameters. Such end-of-line comments are a good place to cite references from which parameter values have been taken.

3.2 Headers

Lines beginning with a bang (!) are called headers (e.g. lines 1–3). Headers are a type of comment designed primarily to store metadata about the reaction network (title, author, date, description etc.). Headers may appear anywhere in the .rxn file (but it is intended that they appear only at the top). All headers will be collected together as comments at the top of the .arb output file. A line showing the date and time at which the .arb output file was generated and the version of `rxntoarb` used will be appended to the header automatically.

3.3 Species and region names

Chemical species have the general form `speciesname@region`. The `speciesname` describes the chemical identity of the species; the `@region` specifies the arb region on/in which that species is present. (Note that `@` is not part of the region name, but rather acts to separate the species and region names.) Chemical species are declared by including them in a reaction (see § 3.5).

Species and region names may not contain any of the characters `<>#` (restrictions inherited from *arb*). `rxntoarb` also requires that species and region names not contain any of the characters `,;@` or any of the substitution sequences or extended syntax used in template files (see § 4). Colons (:) within species names are intended to separate the components of a complex (e.g. `TFPI:Xa` in line 14 is a complex of `TFPI` and `Xa`), but this interpretation only comes into play in calculating the molecular weight of complexes (if molecular weights of the components are given). The use of *arb*-style angle bracket delimiters (`<>`) is unnecessary (but optional) for species names, and necessary for region names only if they contain characters other than letters, digits, and underscores.

Two lazy region identifiers, `@s` and `@v`, are predefined (see § 3.4.1). These identifiers interact with `surface_region` and `volume_region` statements. As such, `s` and `v` should not be used as user-defined region names.

When `rxntoarb` is called with the `-n` flag, ODE mode is activated. In this mode, all species and reactions are assumed to have no spatial dependence and hence region names effectively become part of the species name.

It is permissible to omit `@region` identifiers throughout the `.rxn` file under two circumstances:

- No `surface_region` or `volume_region` statements appear. In this case all species and reactions will be assumed to have no spatial dependence, i.e. ODE mode will be assumed.
- Exactly one surface or volume region is specified by either a `surface_region` or `volume_region` statement. In this case all species and reactions will be assumed to be located on/in that region.

3.4 `surface_region` and `volume_region` statements

The `surface_region` and `volume_region` keywords (both optionally plural) are used to specify *arb* regions on/in which chemical species are present and reactions occur. The keywords should be followed by a comma-separated list of region names (conforming to the naming rules stipulated in § 3.3). If `rxntoarb` is called without the `-n` flag (PDE mode), surface regions will be assumed to be two-dimensional and volume regions will be assumed to be three-dimensional. If `rxntoarb` is called with the `-n` flag (ODE mode), then the dimensionality of regions is ignored.

`surface_region` and `volume_region` statements apply to all reactions subsequently defined in the `.rxn` file, unless and until overridden by another `surface_region` or `volume_region` statement.

3.4.1 Lazy region identifiers `@s` and `@v`

The lazy region identifiers `@s` and `@v` are shorthand references to the surface and volume regions specified by the previous `surface_region` and `volume_region` statements. For example, the `surface_regions` statement in line 20 specifies two surface regions, `TF_patch` and `TM_patch`. References to `@s` in lines 21–24 will therefore be interpreted by `rxntoarb` as references to both `TF_patch` and `TM_patch`, whilst references to `@v` will be interpreted as references to the volume region `domain` specified by the `volume_region` statement in line 6. In line 27, the new `surface_region` statement overrides the previous one so that references to `@s` in lines 28–31 will be interpreted as references to `TF_patch` only.

3.5 Reactions

Four types of reactions are recognised, as described below. Reactions may be entered in any order in the `.rxn` file, subject to the condition that the lazy region identifiers `@s` and `@v` will take their value(s) from the previous `surface_region` and `volume_region` statements.

3.5.1 Irreversible reactions

Irreversible reactions take the form

```
reactant_list -> product_list.
```

`reactant_list` consists of one or more species, separated by spaces and the character `+` (the spaces are mandatory so that `+` can also be used in species names, e.g. for ions). Stoichiometric coefficients (if any) should immediately precede the species name, optionally separated from it by a space or the `.` or `*` characters. `product_list` is similar except that it may be empty, in which case the products of the reaction are discarded (not tracked in the *arb* simulation). For readability, spaces around the reaction arrow `->` are recommended (but not mandatory). The kinetic parameter `k` describing the rate of the reaction must either be specified or inherited (see § 3.6). Examples of irreversible reactions are shown in lines 11–13.

3.5.2 Reversible reactions

Reversible reactions take the form

```
reactant_list <=> product_list.
```

`product_list` must not be empty for a reversible reaction. The kinetic parameters `ka` and `kd` describing the rates of the forward and reverse reactions must either be specified or inherited (see § 3.6). Examples of reversible reactions are shown in lines 23–24.

3.5.3 Two-step reactions

Two-step reactions consist of a reversible reaction followed by an irreversible reaction:

```
reactant_list <=> intermediate_list -> product_list.
```

`intermediate_list` must not be empty but `product_list` may be empty (which implies that the products are not tracked in the *arb* simulation). The kinetic parameters `ka`, `kd`, and `k` describing the rates of the forward and reverse reactions in the first step and the irreversible reaction in the second step must either be specified or inherited (see § 3.6).

3.5.4 Michaelis–Menten reactions

Michaelis–Menten reactions take the form

```
substrate {enzyme}-> product_list.
```

Exactly one `substrate` and one `enzyme` must be specified. `product_list` may be empty (which implies that the products are not tracked in the *arb* simulation). The Michaelis–Menten parameters `KM` and `kcat` must either be specified or inherited (see § 3.6). Examples of Michaelis–Menten reactions are shown in lines 7–10.

3.6 Kinetic parameters

Kinetic parameters are entered as a comma-separated list on the same line as the reaction they describe, separated from the reaction itself by a semicolon (;). Kinetic parameters may be entered in any order. The following parameter names are recognised:

- `k`,
- `ka` (synonyms `kf` and `kon`),
- `kd` (synonyms `kr` and `koff`),
- `KM` (synonym `Km`),
- `kcat`.

Kinetic parameter values may be entered in two forms. In the first form (used throughout the example `.rxn` file), a numerical value together with appropriate physical units are specified. The numerical value and each individual unit should be separated by spaces. SI prefixes immediately precede the unit name; `u` is used for the ‘micro’ prefix. Exponents on units should immediately follow the unit name, optionally preceded by the `^` character. More information on unit entry is given in the accompanying documentation for the `convert.units` program, which provides a command-line interface to the `Units` module that accompanies `rxntoarb`. Where necessary, `rxntoarb` performs unit conversions so that parameter values in the `.arb` output file are expressed in SI base units. Unit consistency checking is also performed with warning messages issued if reaction rates have inconsistent or unexpected units. (Such warnings are produced by lines 23–24.)

Alternatively, kinetic parameter values may be specified as strings, delimited by either single or double quotes. Strings may contain arbitrary *arb* code. This feature is useful in conjunction with aliases (see § 3.7) or for expressing one kinetic parameter in terms of others. The downside is that unit consistency checking is not possible.

3.6.1 Inheritance

Sometimes the same kinetic parameters will apply to a group of closely related reactions. Kinetic parameters from one ‘parent’ reaction can be assigned to any number of ‘child’ reactions by placing the children directly below the parent and indenting them by at least one space. An example of inheritance is shown in lines 15–16: the reaction in line 16 inherits the kinetic parameters from line 15. The advantage of inheritance is that kinetic parameters for the group of reactions need only be entered in one place.

3.7 Aliases

Any non-indented reaction may be preceded by an alias, which runs from the beginning of the line until a colon (:) followed by a (mandatory) space is encountered. Aliases serve two purposes. Firstly, they may be used to provide descriptive names for reactions. Secondly, they provide a convenient way to refer to kinetic parameters defined elsewhere

in the .rxn file (without requiring the user to be familiar with the parameter naming conventions used internally by `rxntoarb`). For example, lines 15–16 could alternatively be written

```
15 Va_inactivation: Va@v {APC@v}-> ; kcat=0.58 s-1, KM=0.025 uM
16 VIIIa@v {APC@v}-> ; kcat="<kcat_Va_inactivation>", KM="<KM_Va_inactivation>"
```

In this case, the values of `kcat` and `KM` specified in line 15 (alias `Va_inactivation`) have been assigned to the reaction in line 16.

3.8 initial_species statement

The `initial_species` statement is used to declare which species are present initially in the *arb* simulation. This information is used by `rxntoarb` to estimate the magnitudes of the concentrations of all species based on the concentrations of the precursor species from which they are derived.

The `initial_species` statement takes a comma-separated list of species together with their associated regions, as shown in line 39. If multiple species on/in the same region are present initially, the array notation shown in line 38 may be used. (Note that lines 38 and 39 have been separated only for aesthetic reasons; they could be combined into a single line.) Any species appearing in an `initial_species` statement but otherwise not appearing in the .rxn file will be ignored.

See § 2.4 for details on how the initial concentrations must be specified in the main .arb file.

3.9 include_only and exclude statements

The `include_only` and `exclude` statements are used to selectively include or exclude regions or reactions in the .arb output file. Both statements must be followed by a valid Ruby regular expression. (The Regexp option `i`, which enables case-insensitive matching, is the only option allowed by `rxntoarb`.) For the statement to operate on a region, the region name must be preceded by the `@` character. If the statement is to operate on a combination of regions and other text (by means of the Regexp alternation operator `|`), then the first alternative must be a region (preceded by `@`). For example, if it is desired to run the example in Listing 1 without the `TM_patch` region, then the following statement should be issued near the top of the .rxn file:

```
exclude /@TM_patch|APC/
```

This statement excludes all reactions occurring on the `TM_patch` as well as those reactions containing the species `APC` (which is not present initially and is only produced on the `TM_patch`).

`include_only` and `exclude` statements apply to all lines subsequently defined in the .rxn file, unless and until overridden by another `include_only` or `exclude` statement. End-of-line comments are included in the matching process, effectively enabling the user to define their own keywords to select/deselect subsets of reactions.

4 Template file syntax

`rxntoarb` template files (called `rxntoarbrc` by default) store *arb* code that needs to be included for each species. A range of substitutions are performed on the original template file to generate unique code for each species. Template files contain extended (non-*arb*) syntax that is handled by `rxntoarb` as detailed below.

4.1 Substitution sequences

`rxntoarb` performs substitutions on the following character sequences (all delimited by forward slashes), which must therefore not appear in any species or region names:

- `/c/` (concentration of the species),
- `/species/` (species name),
- `/region/` (region name),
- `/source_region/` (region name),
- `/CENTRING/` (NONE|FACE|CELL depending on species location),
- `/centring/` (none|face|cell depending on species location),
- `/units/` (''|mol m-2|mol m-3 depending on species location),
- `/MW/` (molecular weight of the species),
- `/associatedfaces(region)/` (region consisting of faces associated with the region),
- `/associatedcells(region)/` (region consisting of cells associated with the region),
- `/domainof(region)/` (region consisting of cells in domain of the region).

Additionally, all *arb* statements of the form `ON <region>` are deleted in ODE mode.

4.2 `if_rxn` statement

`if_rxn` statements are used to conditionally include *arb* code depending on whether a species is located on a surface, in a volume, or has no location (ODE mode). `if_rxn` statements take the form

```
if_rxn(location=region_list){if_clause}{else_clause}.
```

The term in parentheses determines the scope of the `if_rxn` statement. `location` must be one of `none`, `surface`, or `volume`. The optional `region_list` limits the scope of the

`if_rxn` statement to only the named regions, and should be a comma-separated list of region names (*arb*-style angle bracket delimiters are optional). The code contained in the mandatory `if_clause` is retained for all species defined within the scope of the `if_rxn` statement and deleted otherwise. The code contained in the optional `else_clause` is retained for all species outside the scope of the `if_rxn` statement and deleted otherwise. Both `if_clause` and `else_clause` may span multiple lines. Nested brace pairs within clauses are allowed but nested `if_rxn` statements are not permitted.

4.3 `each` blocks

`each` blocks are used to apply a particular boundary condition to multiple surface regions bounding a volume region, and hence must reside within `if_rxn(volume)` statements. `each` blocks take the Ruby-inspired form

```
[region_list].each { |loopvar| expression }.
```

The mandatory `region_list`, which specifies the volume-region-bounding surface regions for which the `expression` in the block will be retained, should be a comma-separated list of region names (*arb*-style angle bracket delimiters are optional). Each of these region names (stripped of angle brackets, if any) will be passed to the block, where they are accessed through the (user-named) variable `|loopvar|` (vertical bars mandatory). In `region_list`, the shorthand `*` may be used to represent all volume-region-bounding surface regions on which reactions are occurring.

`each` blocks may span multiple lines. Nested brace pairs within `each` blocks are allowed but nested `each` blocks are not permitted.

5 Future development ideas

No guarantees, just some ideas! Contact the author if there is a capability you think `rxntoarb` should have.

- Top priority: support conversion between `.rxn` format and `SBML` format. Potentially `SBML` could completely replace the homegrown `.rxn` format in future.
- Allow subsets of reactions to occur at different ‘runlevels’, e.g. to simulate sequential reaction steps. Would be handy if pre-equilibration steps could be run through *arb* automatically. (Currently this can be achieved by making use of `include_only` and writing a wrapper script.)
- \LaTeX output of reactions. (Could defer to `COPASI` or `SBML2LaTeX` if `SBML` support implemented.)
- Graphical output of reaction network diagrams. (Could defer to `COPASI` or `Cytoscape` if `SBML` support implemented.)

6 Copyright and licence

`rxntoarb` source code and documentation © 2016–2017 Christian Biscombe.

`rxntoarb` is contributed to *arb* finite volume solver (in which copyright is held by Dalton Harvie) under the same licence terms as that project. At the time of writing, *arb* is released under the terms of the GNU General Public License (version 3) as published by the Free Software Foundation.